



并行与分布式计算

Parallel & Distributed Computing

陈鹏飞
数据科学与计算机学院





Lecture 5 — OpenMP Programming

Pengfei Chen

School of Data and Computer Science



Outline:



OpenMP Overview



Correctness



Task



Performance



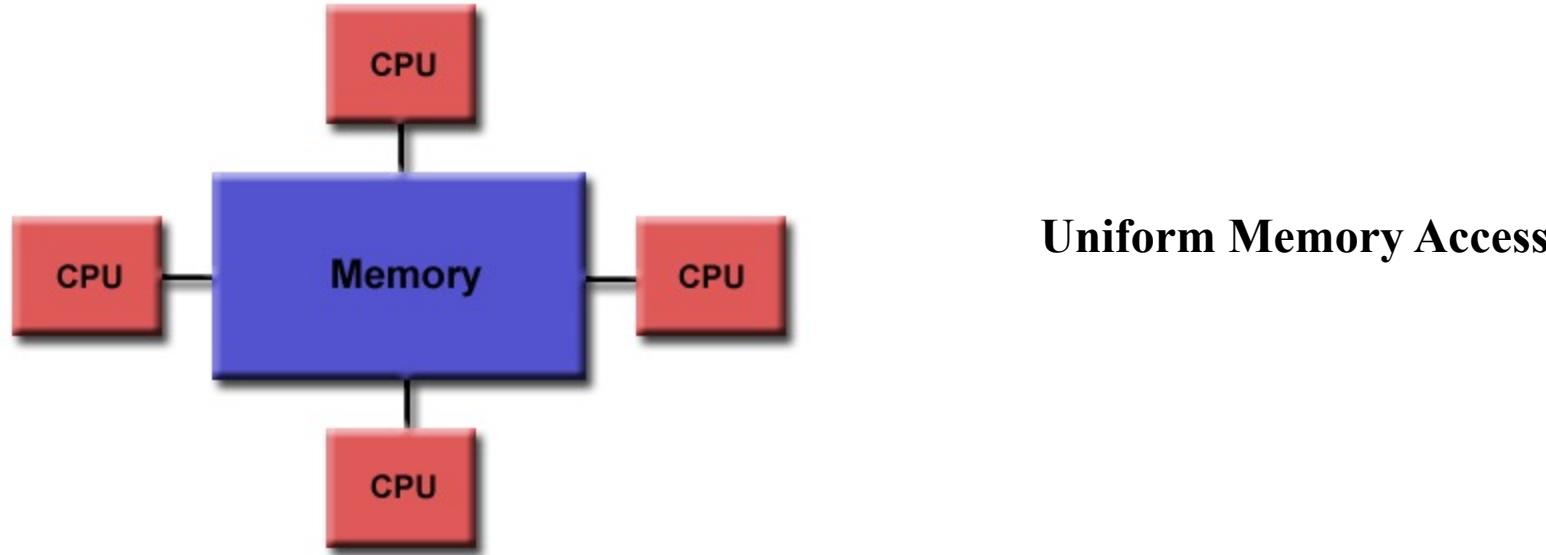


OpenMP programming

Overview

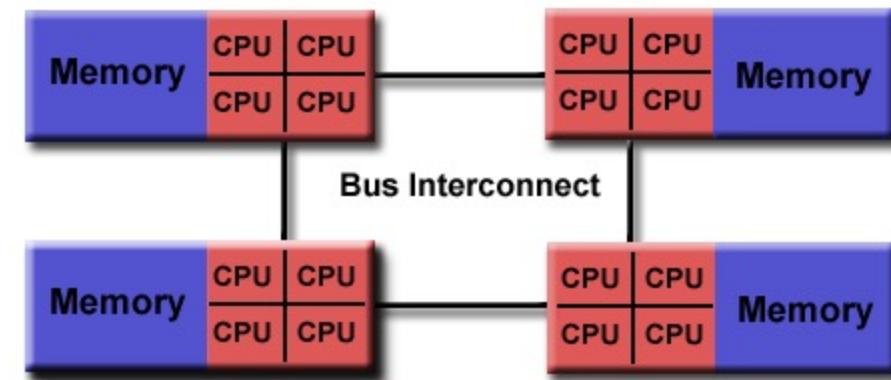


Architecture for Shared Memory Model



Uniform Memory Access

Non-uniform Memory Access





Thread Base Parallelism

- OpenMP programs accomplish parallelism exclusively (仅仅) through the use of threads
- A thread of execution is the smallest unit of processing that can be scheduled by an operating system
 - The idea of a subroutine that can be scheduled to run autonomously might help explain what a thread is
- Threads exist within the resources of a single process
 - Without the process, they cease (停止) to exist
- Typically, the number of threads match the number of machine processors/cores
 - However, the actual use of threads is up to the application



Explicit Parallelism

- OpenMP is an explicit (not automatic) programming model, offering the programmer full control over parallelization
- Parallelization can be as simple as taking a serial program and inserting compiler directives....
- Or as complex as inserting subroutines to set multiple levels of parallelism, locks and even nested locks



What is OpenMP?

- An abbreviation for
 - Short version
 - **Open Multi-Processing** (开放多处理过程)
 - Long version
 - **Open specifications for Multi-Processing via collaborative work between interested parties from the hardware and software industry, government and academia**



OpenMP Overview

C\$OMP FLUSH

#pragma omp critical

C\$OMP THREADPRIVATE (/ABC/)

CALL OMP SET NUM THREADS (10)

OpenMP: An API for Writing Multithreaded Applications

- A set of compiler directives and library routines for parallel application programmers
- Greatly simplifies writing multi-threaded (MT) programs in Fortran, C and C++
- Standardizes last 20 years of SMP practice

C\$OMP PARALLEL COPYIN (/blk/)

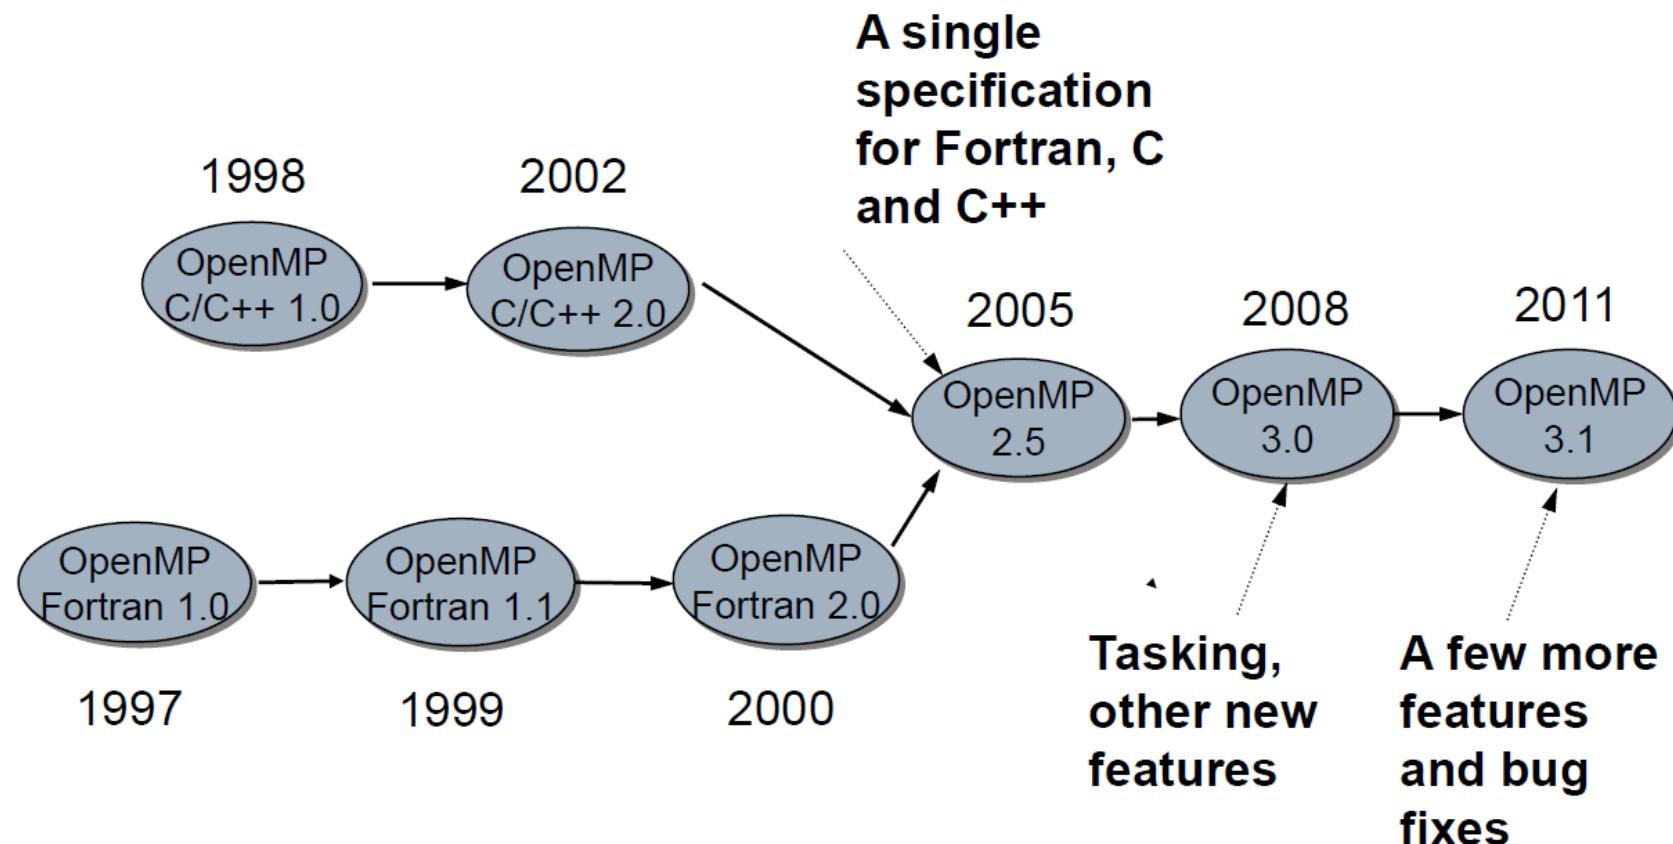
C\$OMP DO lastprivate (XX)

Nthrds = OMP_GET_NUM_PROCS()

omp_set_lock(lck)



OpenMP Release History





OpenMP Overview: How do Threads Interact?

- OpenMP is a multi-threading, shared address model
 - Threads communicate by sharing variables
- Unintended sharing of data causes race conditions
 - Race condition: when the program's outcome changes as the threads are scheduled differently
- To control race conditions
 - Use synchronization to protect data conflicts
- Synchronization is expensive
 - Change how data is accessed to minimize the need for synchronization



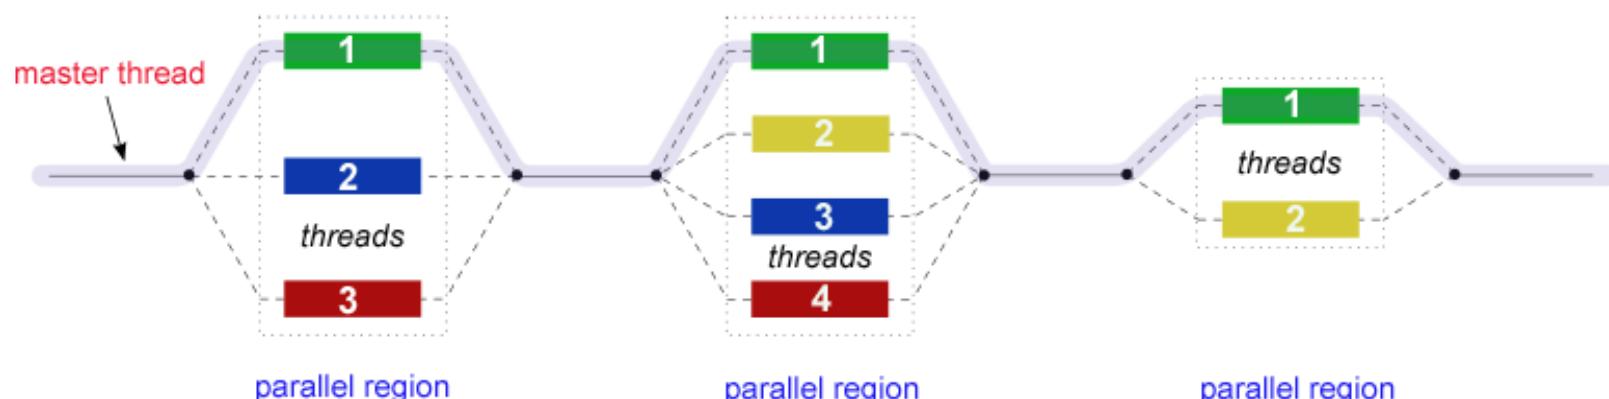
Fork-Join Model

➤ FORK

- The master thread then creates (or awakens) a team of parallel threads.

➤ JOIN

- When the team threads complete the statements in the parallel region construct, they synchronize and terminate, leaving only the master thread.





Relating Fork/Join to Code



Sequential code

Parallel code

Sequential code

Parallel code

Sequential code



OpenMP Components

- Three API components
 - Compiler directives
 - Runtime library routines
 - Environment variables



Syntax of Compiler Directives (指令)

- *pragma*: a C/C++ compiler directive (编译开关)
 - (other compiler directives: #include, #define, ...)
 - Stands for “pragmatic information (附注信息)”
 - A way for the programmer to communicate with the compiler
 - Pragmas are handled by the preprocessor
 - Compilers are free to ignore pragmas
- All OpenMP pragmas have the syntax:
 - #pragma omp <directive-name> [clause, ...]
- Pragmas appear immediately **before** relevant construct



Hello World

```
#include <omp.h>
#include <stdio.h>
#include <stdlib.h>
int main (int argc, char *argv[]) {
    int nthreads, tid;
    /* Fork a team of threads giving them their own copies of variables */
#pragma omp parallel private(nthreads, tid)
{
    /* Obtain thread number */
    tid = omp_get_thread_num();
    /* Only master thread does this */
    if (tid == 0) {
        nthreads = omp_get_num_threads();
        printf("Number of threads = %d\n", nthreads);
    }
    printf("Hello World from thread = %d\n", tid);
} /* All threads join master thread and disband */
}
```



Output – Non-deterministic!

```
[cong@lnxsrvg1 ~/cs133_examples]$ ./a.out
Hello World from thread = 5
Hello World from thread = 7
Number of threads = 16
Hello World from thread = 0
Hello World from thread = 8
Hello World from thread = 3
Hello World from thread = 1
Hello World from thread = 6
Hello World from thread = 15
Hello World from thread = 14
Hello World from thread = 10
Hello World from thread = 2
Hello World from thread = 9
Hello World from thread = 11
Hello World from thread = 13
Hello World from thread = 12
Hello World from thread = 4
[cong@lnxsrvg1 ~/cs133_examples]$ ./a.out
Hello World from thread = 6
Hello World from thread = 10
Hello World from thread = 8
Number of threads = 16
Hello World from thread = 0
Hello World from thread = 5
Hello World from thread = 14
Hello World from thread = 15
Hello World from thread = 13
Hello World from thread = 12
Hello World from thread = 4
Hello World from thread = 3
Hello World from thread = 9
Hello World from thread = 2
Hello World from thread = 7
Hello World from thread = 1
Hello World from thread = 14
```



OpenMP Compilers Perform the Translations to Threads (e.g. pthreads)

```
int a, b;
main() {
    // serial segment
    #pragma omp parallel num_threads (8) private (a) shared (b)
    {
        // parallel segment
    }
    // rest of serial segment
}
```

Sample OpenMP program

```
int a, b;
main() {
    // serial segment
    Code inserted by the OpenMP compiler
    for (i = 0; i < 8; i++)
        pthread_create (....., internal_thread_fn_name, ...);
    for (i = 0; i < 8; i++)
        pthread_join (.....);
    // rest of serial segment
}

void *internal_thread_fn_name (void *packaged_argument) {
    int a;
    // parallel segment
}
```

Corresponding Pthreads translation



Matching Threads with CPUs

- ◆ Function `omp_get_num_procs` returns the number of physical processors available to the parallel program

```
int omp_get_num_procs(void);
```

- ◆ Function `omp_set_num_threads` allow you to set the number of threads that should be active in parallel sections of code

```
void omp_set_num_threads(int t);
```

- The function can be called with different arguments at different points in the program



Pragma: parallel for

◆ The compiler directive

#pragma omp parallel for

tells the compiler that the *for* loop which immediately follows can be executed in parallel

- The number of loop iterations must be computable at run time before loop executes
- Loop must not contain a *break*, *return*, or *exit*
- Loop must not contain a *goto* to a label outside loop



Example: parallel for

```
int a[1000], b[1000], s[1000];
```

```
...
```

```
#pragma omp parallel for
for (i = 0; i < 1000; i++)
    s[i] = a[i] + b[i];
```

- ◆ Threads are assigned an independent set of iterations
- ◆ Threads must wait at the end of construct



Which Loop to Make Parallel?

```
int main() {  
    int i, j, k;  
    float **a, **b;  
    ... // initialize a[][][], b[][] as the 1-hop distance matrix  
    for (k = 0; k < N; k++) {  
        for (i = 0; i < N; i++)  
            for (j = 0; j < N; j++)  
                a[i][j] = min(a[i][j], b[i][k] + b[k][j]);  
    }  
    ... // copy a[][] to b[][]
```



Which Loop to Make Parallel?

```
int main() {  
    int i, j, k;  
    float **a, **b;  
    ... // initialize b[][] as the 1-hop distance matrix  
    for (k = 0; k < N; k++)          // Loop-carried dependences  
        for (i = 0; i < N; i++)      // Can execute in parallel  
            for (j = 0; j < N; j++)  // Can execute in parallel  
                a[i][j] = min(a[i][j], b[i][k] + b[k][j]);  
    ... // copy a[][] to b[][]
```



Minimizing Threading Overhead

- There is a fork/join for every instance of

```
#pragma omp parallel for
```

```
for (...) {
```

```
...
```

```
}
```

- Since fork/join is a source of overhead, we want to maximize the amount of work done for each fork/join; i.e., **the grain size**
- Hence we choose to make the middle loop parallel



Almost Right, but Not Quite

```
int main() {  
    int i, j, k;  
    float **a, **b;  
    ... // initialize b[][] as the 1-hop distance matrix  
    for (k = 0; k < N; k++) {  
        #pragma omp parallel for  
        for (i = 0; i < N; i++)  
            for (j = 0; j < N; j++)  
                a[i][j] = min(a[i][j], b[i][k] + b[k][j]);  
    ... // copy a[][] to b[][]  
}
```

Problem: j is a shared variable



Clause: private

- Clause
 - An optional, additional component to a pragma
- Private clause: directs compiler to make one or more variables

private

```
#pragma omp ... private (<variable list>)
```



Problem Solved with private Clause

```
int main() {  
    int i, j, k;  
    float **a , **b;  
    ... // initialize b[][] as the 1-hop distance matrix  
    for (k = 0; k < N; k++) {  
        #pragma omp parallel for private (j)  
        for (i = 0; i < N; i++)  
            for (j = 0; j < N; j++)  
                a[i][j] = min(a[i][j], b[i][k] + b[k][j]);  
    ... // copy a[][] to b[][]  
}
```

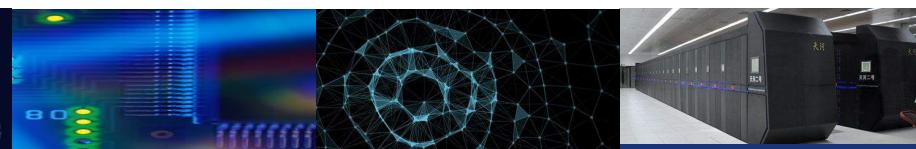
Tell compiler to make listed variables private



Another Example

```
int i;  
float *a, *b, *c, tmp;  
  
...  
  
for (i = 0; i < N; i++) {  
    tmp = a[i] / b[i];  
    c[i] = tmp * tmp;  
}
```

**Loop is perfectly parallelizable except for shared
variable *tmp***



Solution

```
int i;  
  
float *a, *b, *c, tmp;  
  
...  
  
#pragma omp parallel for private (tmp)  
  
for (i = 0; i < N; i++) {  
  
    tmp = a[i] / b[i];  
  
    c[i] = tmp * tmp;  
  
}
```



More About Private Variables (私有变量)

- Each thread has its own copy of the private variables
- If j is declared private, then inside the *for* loop no thread can access the “other” j (the j in shared memory)
- No thread can use a previously defined value of j
- No thread can assign a new value to the shared j
- Private variables are **undefined** at loop entry and loop exit, reducing execution time



Clause: firstprivate

- The ***firstprivate* clause tells the compiler that the private variable should inherit the value of the shared variable upon loop entry**
- The value is assigned once per thread, not once per loop iteration



Example: firstprivate

```
a[0] = 0.0;  
for (i = 1; i < N; i++)  
    a[i] = alpha(i, a[i-1]);
```

```
#pragma omp parallel for firstprivate (a)  
for (i = 0; i < N; i++)  
    b[i] = beta(i, a[i]);  
    a[i] = gamma(i);  
    c[i] = delta(a[i], b[i]);  
}
```



Clause: firstprivate

➤ pragma omp ... firstprivate(x)

□ ***x is a fundamental data type***

- Private x is directly copied from the shared x

□ ***x is an array***

- Copy the data with $\text{sizeof}(x)$ to the private memory

□ ***x is a pointer***

- Private x points to the same location as the shared x

□ ***x is a class instance***

- Copy constructor is called to create the private x



Clause: firstprivate

- pragma omp ... firstprivate(x)
 - ***x is a fundamental data type***
 - Private *x* is directly copied from the shared *x*
 - ***x is an array***
 - Copy the data with *sizeof(x)* to the private memory
 - ***x is a pointer***
 - Private *x* points to the same location as the shared *x*
 - ***x is a class instance***
 - Copy constructor is called to create the private *x*



Clause: lastprivate

- The ***lastprivate*** clause tells the compiler that the value of the private variable after the ***sequentially last*** loop iteration should be assigned to the shared variable upon loop exit
 - In other words, when the thread responsible for the ***sequentially last*** loop iteration exits the loop, its copy of the private variable is copied back to the shared variable



Example: lastprivate

```
#pragma omp parallel for lastprivate (x)
```

```
for (i = 0; i < N; i++)
```

```
    x = foo(i);
```

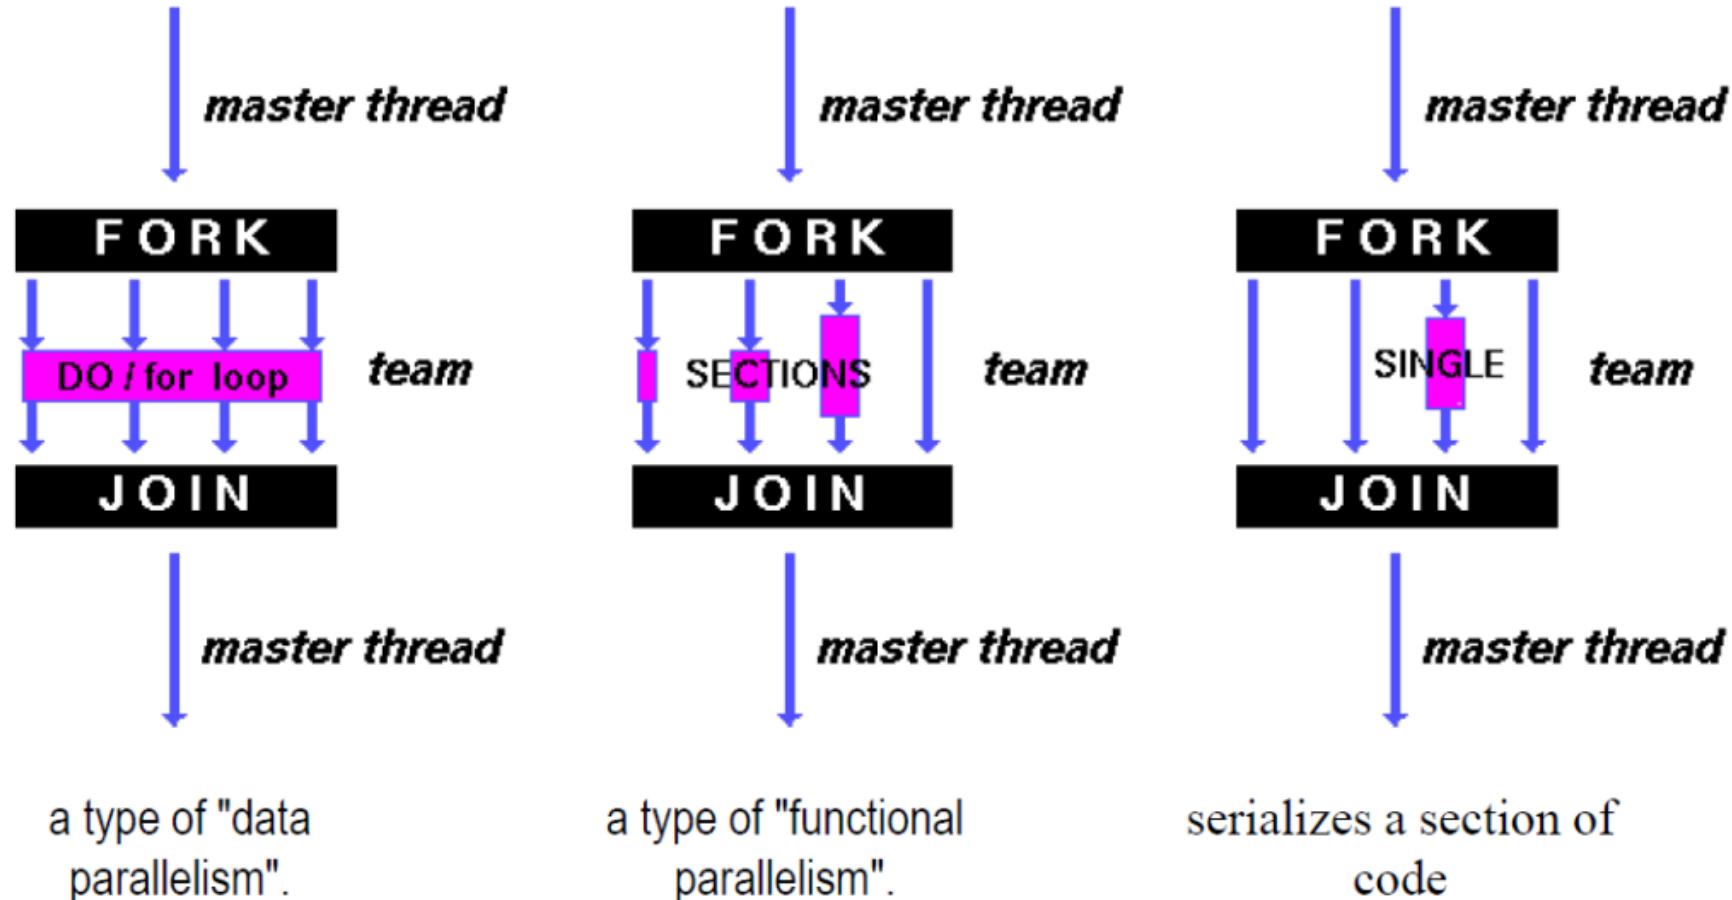
```
    y[i] = bar(i, x);
```

```
}
```

```
last_x = x; // == foo(N-1)
```



Work-Sharing Constructs





Pragma: parallel

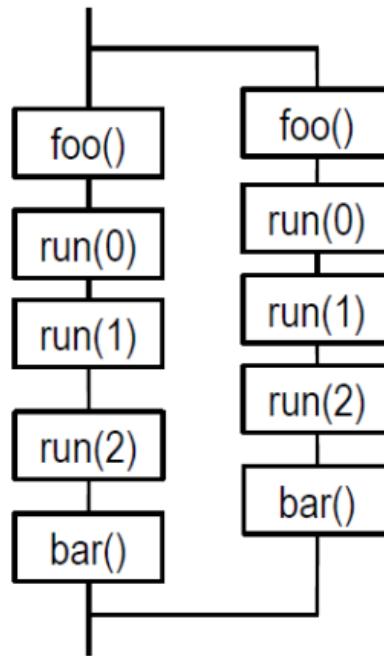
- In the effort to increase grain size, sometimes the code that should be executed in parallel goes beyond a single *for* loop
 - The *parallel* pragma is used when a block of code should be executed in parallel
 - SPMD-style programming
 - Single program, multiple data



Pragma parallel and Pragma for

◆ #pragma omp for

- Used inside a block of code already marked by *parallel*
- Distribute the iterations to the active threads



```
#pragma omp parallel \
num_threads(2)
{
    foo();
    for (int i = 0; i < 3; i++)
        run(i);
    bar();
}
```

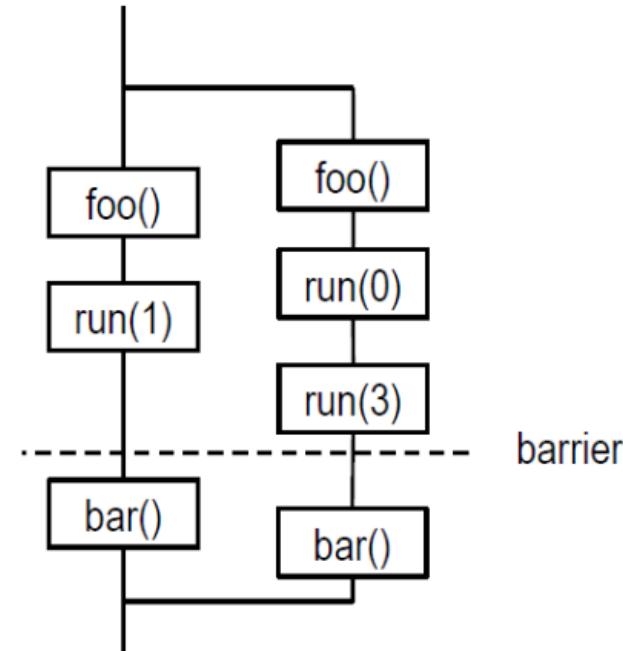


Pragma parallel and Pragma for

◆ #pragma omp for

- Used inside a block of code already marked by *parallel*
- Distribute the iterations to the active threads

```
#pragma omp parallel \
num_threads(2)
{
    foo();
    #pragma omp for
    for (int i = 0; i < 3; i++)
        run(i);
    bar();
}
```





Pragma: single

- The *single* pragma is used inside a parallel block of code
 - It tells the compiler that only a single thread should execute the statement or block of code immediately following
 - May be useful when dealing with sections of code that are not thread safe (such as I/O)
 - Threads in the team that do not execute the single directive, wait at the end of the enclosed code block, unless a nowait clause is specified.



Example: master and single nowait

```
tid = omp_get_thread_num();
if (tid == 0) {
    nthreads = omp_get_num_threads();
    printf("Number of threads = %d\n", nthreads);
}
```

=

```
#pragma omp master
{
    nthreads = omp_get_num_threads();
    printf("Number of threads = %d\n", nthreads);
}
```

≈

```
#pragma omp single nowait
{
    nthreads = omp_get_num_threads();
    printf("Number of threads = %d\n", nthreads);
}
```



Pragma: sections and section

- Directive sections specifies that the enclosed section(s) of code are to be divided among the threads in the team.
- Independent section directives are nested within a sections directive.
 - Each section is executed once by a thread in the team.
 - Different sections may be executed by different threads.
 - It is possible for a thread to execute more than one section if it is quick enough and the implementation permits such.



Example: sections and section

```
#pragma omp parallel shared(a,b,c,d) private(i)
{
#pragma omp sections
{
#pragma omp section
{
    for (i=0; i<N; i++)
        c[i] = a[i] + b[i];
}
#pragma omp section
{
    for (i=0; i<N; i++)
        d[i] = a[i] * b[i];
}
} /* end of sections */
} /* end of parallel section */
```



Clause: reduction (归并)

- ◆ Reductions are so common that OpenMP provides a reduction clause for the *parallel, for, and sections*

```
#pragma omp ... reduction (op : list)
```

- A **private** copy of each list variable is created and initialized depending on the *op*
 - The identity value *op* (e.g., 0 for addition)
- These copies are **updated locally** by threads
- At end of construct, local copies are combined through *op* into a single value and combine the value in the original **shared** variable



C/C++ Reduction Operation

结合率

- ◆ Reduction with an **associative** binary operator \oplus

$$a_1 \oplus a_2 \oplus a_3 \oplus \dots \oplus a_n$$

- ◆ A range of associative and commutative operators can be used with reduction
- ◆ Initial values are the ones that make sense

Operator	Initial Value
+	0
*	1
-	0
\wedge	0

Operator	Initial Value
$\&$	~ 0
$ $	0
$\&\&$	1
$\ $	0



Strengths and Weaknesses of OpenMP

➤ Strengths

- Incremental parallelization & sequential equivalence
- Well-suited for domain decompositions
- Available on *nix and Windows

➤ Weaknesses

- Not well-tailored for functional decompositions
- Compilers do not have to check for such errors as deadlocks and race conditions



Homework

- Consider a sparse matrix stored in the compressed row format (you may find a description of this format on the web or any suitable text on sparse linear algebra). Write an OpenMP program for computing the product of this matrix with a vector. Download sample matrices from the Matrix Market (<http://math.nist.gov/MatrixMarket/>) and test the performance of your implementation as a function of matrix size and number of threads.



A visual repository of test data for use in comparative studies of algorithms for numerical linear algebra, featuring nearly 500 sparse matrices from a variety of applications, as well as matrix generation tools and services.

Browse by collection	Search by matrix properties	Background Welcome
by matrix name	by application area	What's New
by generator name	by contributor	What's Coming
		Credits

1138 BUS: Power systems admittance matrices Power system networks

from set [PSADMIT](#), from the [Harwell-Boeing Collection](#)

[\[Download\]](#) [\[Visualizations\]](#) [\[Matrix Statistics\]](#) [\[Set Information\]](#)

Download as

- Compressed [MatrixMarket format](#) file: [1138_bus.mtx.gz](#) (21322 bytes)
- Compressed [Harwell-Boeing format](#) file: [1138_bus.rsa.gz](#) (19648 bytes)

Help: My browser can't read the compressed data files. [What now?](#)

- Implement a producer-consumer framework in OpenMP using sections to create a single producer task and a single consumer task. Ensure appropriate synchronization using locks. Test your program for a varying number of producers and consumers.



中山大学 计算机学院（软件学院）

SUN YAT-SEN UNIVERSITY

SCHOOL OF COMPUTER SCIENCE AND ENGINEERING



Thank You !